

Detection of SQL Injection Attacks by giving apriori to Q-Learning Agents

Tejas Sheth

Department of Computer Engineering
A. P. Shah Institute of Technology
Thane (M.H.), India, 400615
tejas.sheth04@gmail.com

Janhavi Anap

Department of Computer Engineering
A. P. Shah Institute of Technology
Thane (M.H.), India, 400615
anapjanhavi@gmail.com

Het Patel

Department of Computer Engineering
A. P. Shah Institute of Technology
Thane (M.H.), India, 400615
hetpokar@gmail.com

Nidhi Singh

Department of Computer Engineering
A. P. Shah Institute of Technology
Thane (M.H.), India, 400615
singhnidhi2002@gmail.com

Prof. Ramya R B

Department of Computer Engineering
A. P. Shah Institute of Technology
Thane (M.H.), India, 400615
ramyarb@apsit.edu.in

Abstract—Software developers may have created the SQL Injection vulnerability accidentally, or a hacker may have purposefully used it to target vulnerable data. With the recent surge in information, there is an innate quest to safeguard this information from falling into the wrong hand leading to data theft, leak of personal data or loss of property. With relational databases like MySQL being the most popular, it allows users to extract any available information without any significant knowledge of databases. With vast information stored in databases warrants attacker's attention, potentially risking critical confidential information. The premature detection of SQL Injection Attacks will be very helpful in preventing any malicious attempt by an attacker. In this study, we examine the outcomes of algorithms for reinforcement learning, such as Q-Learning, using a dataset made up of probable SQL Injection queries. We intend to provide a Reinforcement Learning solution to minimise the potential threat posed by SQL Injection and give apriori to the model to learn to detect a SQL attack and prevent any unforeseen mishap more quickly and accurately.

Keywords—Machine Learning, Information security, SQL Injection, SQL detection, SQL injection Attacks, Reinforcement Learning, Q-Learning, Vulnerability detection, Apriori

I. INTRODUCTION

Cybersecurity is the method of defending systems and programs from digital strikes. These strikes are usually designed for retrieving, manipulating or destroying sensitive data or interrupting standard business approaches. Cybersecurity aims to reduce the risk of cyber-attacks and to protect against the illegal manipulation of networks and systems. Implementing robust and reliable measures is challenging due to the ever-evolving innovative attackers.

Structured Query Language Injection Attacks (SQLiA) are one of the most dangerous security flaws that are frequently employed by hackers, according to the Top 10 report from the Open Web Application Security Project (OWASP) for 2022 [1]. By inserting malicious data into a database request, hackers can get access to a database without authorization by using SQL Injection (SQLi).

Data held in relational databases can be created, retrieved, updated, and deleted using the computer programming language SQL. All relational database management systems (RDBMS), including MySQL, employ SQL, a standard database language, extensively. SQL has the advantage of allowing users to access, define, describe, and alter data. Additionally, it permits users to make databases, drop and create tables, make views, store procedures, and functions, as well as control permissions for tables, views, and procedures. SQL enables the use of pre-compilers, libraries, and modules with other languages.

The main reason why SQLi keeps traction is due to improper and insecure development architecture. Making use of legacy software introduces security vulnerabilities that might not be detected with modern software. Using authorized and the latest versions of the software are critical to avoiding or minimizing potential security exploits, including SQLiA.

SQLiA mainly occurs through the concatenation of a malicious SQL command at the end of a given input query on the user front end. In recent times, there has been an advancement in the number of dynamic websites, that rely on relational databases to store and manipulate large amounts of data to provide a rich user experience. The SELECT command in conjunction with the WHERE clause is used to retrieve data, which is one of the most popular SQL actions.

A few of the common SQLiA types are Piggy-Backed, Stored Procedures, Union Query, and Alternative Encoding according to Halfond, W. G. et. al. [2] and Wei, K. et. al. [3]. There are some advanced SQLiA types like Blind SQLi where the attackers devise a strategy to circumvent the lack of error notifications, Fast Flux SQLi where the DNS method is used to conceal phishing and malware distribution sites behind a constantly changing network of compromised servers, Compounded SQLi is when the attacker combines two or more attacks that target the website and has improbable repercussions.

The detection of these attack types becomes a very important task keeping in mind the potential data breaches that could happen. The detection techniques include keeping an eye on database issues and reviewing server logs. The most popular method for identifying SQL Injection threats is pattern matching, which combines static analysis and real-time monitoring.

Another approach to detecting SQLiA is by making use of various Machine Learning techniques like supervised and unsupervised, Reinforcement Learning. There are many supervised and unsupervised approaches towards detecting SQLiA.

In this paper, we have used the Reinforcement Learning method towards SQLiA detection. Here, in the training phase, we provide apriori to the Q-Learning model. The training dataset contains both malicious and non-malicious queries. Providing efficient apriori is the most crucial factor in order to obtain higher efficiency. The advantage of Reinforcement Learning techniques over others is that it allows us to explore a variety of SQL queries and the accuracy keeps increasing with the increase in the use of the model. Also, owing to Reinforcement Learning methodologies, the accuracy value of SQL detection is raised and the false positive rate drops.

II. LITERATURE SURVEY

Rai, A. et. al. illustrate the classification and prevention of different SQLi attacks. SQLi is generally classified as In-band SQLi, Inferential SQLi and Out of Bound SQLi [4]. In-Bound SQL injection is further classified as Error-based and Union-based SQLi. Inferential SQLi can be broken down into Boolean-based Blind SQL and Time-based SQL. Defensive techniques that could be used to prevent an SQLi attack include Whitelisting/Blacklisting, prepared statement/ parameterized query, stored procedure, defensive coding practice, taint-based approach, proxy filters, instruction set randomization, low privileges and output Escaping. Different countermeasures work for different SQL Attacks.

Erdodi, L. et. al. simplified the dynamics of SQLi vulnerabilities by projecting the problem as capture-the-flag security and implementing it as an RL problem [5]. Assuming the vulnerability has been recognized, they rely on RL algorithms to automate the process of exploiting SQLi. They implemented the model using two simulations. The first simulation showed that a simple RL agent-based Q-Learning algorithm can successfully develop an effective strategy to resolve the SQLi problem. Through pure trial and error, a tabular Q-Learning algorithm can uncover an effective strategy and achieve performance levels close to the theoretical optimum. However, while the use of a table to store the Q-value function enables a detailed analysis of the agent's learning dynamics, this approach suffers from poor scalability, limiting its practical applicability. Thus, in the second simulation, they sacrificed interpretability to work around the issue of scalability. In the first simulation, a deep Q-Learning agent was used to tackle the issue. It was able to learn an effective approach for the SQLi problem and also offer a solution to the space restrictions caused by the instantiation of an explicit Q-Table.

Verme, M. D. et. al. contemplated the problem of exploiting SQLi flaws and passing it off as a capture-In a "the-flag" scenario, an attacker may fill out a form with strings to obtain a flag token that represents personal data. [6]. The attacker was modelled as an RL agent that maintains interaction with the server to learn an optimal policy leading to the attacker taking advantage of it. The authors did a comparison between two types of agents, one was a simple structured agent that relied on significant apriori knowledge and used high-level actions while the other was a structureless agent that had limited apriori knowledge and generated SQL statements. The comparison demonstrated the viability of creating agents with fewer ad hoc modelling requirements.

Tang, P. et.al. only extracted and classified the URL features [7]. The factors like payload length, keywords and their weights are considered for feature extraction. Using Artificial Neural Network (ANN) models, the URL is categorised as harmful or non-malicious. The method and algorithm used here are multi-layer perceptrons (MLP) and LSTM, both implemented using Pytorch. The trained model is deployed in the ISP system so that abnormal behaviours can be found in the network in real-time. While using the LSTM model for recognition purposes in this approach, there were some drawbacks. These include reduced accuracy, poor model recognition, and increased processing time, which can be problematic for the overall performance of the system.

Niculae, S. et al. measured the performance of multiple fixed-strategy and learning-based agents [8]. They concluded that Q-Learning, with some extra techniques applied and greedy agent initialisation, performed best, surpassing human performance in the given environment.

Hu, Z., Beuran, R., & Tan, Y. suggest an automated penetration testing framework, based on deep learning techniques, particularly Deep Q-Learning networks (DQN) [9]. The authors conducted an experiment in which a given network host was populated with real host and vulnerable data, to determine the optimal attack path, and to provide viable solutions.

RL is fairly successful in tackling and solving games, penetration testing, when distilled as a capture-the-flag (CTF), can be expressed as a game. When it comes to penetration testing, a machine may be limited to learning by trial and error, whereas a human hacker can utilize other methods such as knowledge from alternative sources, deduction, hypothesis testing, and social engineering. Although an RL agent may learn the structure completely without using any models, this may be computationally difficult. Thus according to Zennaro, F. M., & Erdodi, L. injecting some form of elementary apriori knowledge about the structure of the problem may simplify the learning problem [10]. Some basic forms of apriori knowledge are lazy loading, state aggregation and imitation learning, which makes the RL agent much more efficient. The authors categorized CTFs in groups according to the type of vulnerability they instantiate and the type of exploitation that a player is expected to perform. The prototypical classes of CTF problems considered were port scanning and intrusion, server hacking and website hacking. The common RL interface

specified in the OpenAI gym library was used to implement each simulation.

Ghanem M. C., & Chen T. M. propose and evaluates an AI-based pen-testing system which makes use of RL to learn and replicate standard and complicated pen-testing activities [11]. The scope is limited to network infrastructure PT planning and not the entire practice. Moreover, the authors tackle the tricky problem of expertise capturing by allowing the learning module to store and reuse PT policies in a more efficient way.

John, A. contemplated methods incorporating the best qualities of parse tree validation and code conversion techniques [12]. The algorithm parses the user input and checks for vulnerability, and if found, it applies code conversion over that input. Results show few drawbacks of code conversion as applying it to every user input is labour-intensive as well as the database increases. The parse tree validation technique could raise a false alarm if a legitimate user is having blank space in their input. The proposed method proved to provide higher security levels than the individual techniques of code conversion and parse tree validation.

Hanmanthu, B. et. al. illustrates the use of the famous decision tree classification techniques to prevent SQLi attacks [13]. The considered model works by sending different particularly planned attack requests to the proposed SQLi decision tree model, and the final SQLi database is created for using classification data. It uses the satisfied analysis technique for finding the SQLi attack and uses the SQL decision tree. Typically, software developers employ string concatenation to dynamically build SQL statements. However, this approach can be prone to errors and may require manual intervention. The proposed method allows for the creation of multiple queries tailored to meet the different conditions specified by users. This approach eliminates the need for manual intervention and mitigates the risks associated with error-prone coding. The model showed consistency in attack detection and elimination at an average of 82% for all types of attacks.

By fusing a malicious SQL query with the input parameters, SQL Injection introduces a malicious SQL query into a web application. Sadeghian, A. et. al. illustrate how injection attacks are categorized, including tautology inquiries, illegal or illogical questions, union queries, piggy-backed queries, stored procedures, inference, and alternative encoding. [14]. Security researchers have categorized the solutions for SQLi into three major categories: Best code practices, SQLi detection and SQLi runtime prevention. The optimum solution would be writing secure code and among best code practices- parameterized querying is the most secure and efficient technique.

Medhane and M. H. A. S. based their approach on grammar to determine the injection vulnerabilities during software development and SQLi attack based on web applications [15]. By employing SQL queries, the attackers were able to execute their assault, causing a modification in the program's behaviour as the SQL queries were modified, thus disrupting its intended function.

Reinforcement Learning (RL) is known to obtain knowledge

by trial-and-error method and continuous interaction with a dynamic environment. It is characterized by self-improving and online learning, making it one of the intelligent agents (IA) core technologies. In reinforcement learning (RL), the signal provided by the environment serves as a form of evaluation for the quality of the actions taken by the AI. However, this signal does not provide instructions on how to generate the correct action. The basic model of RL as stated in Qiang, W., & Zhongli, Z. includes a state, action and reward system [16], where the IA observes the surroundings and decides how to act to continuously engage with the surroundings to maximise reward value. The ultimate goal of RL is to learn an action strategy. At the core of reinforcement learning technology lies the idea that a system's actions which yield positive rewards from the environment will reinforce its tendency to produce similar actions in the future, thereby creating a positive feedback loop. Conversely, actions that do not produce a positive reward will weaken the system's tendency to generate similar actions. Typical RL method based on the Markov decision-making process (MDP) model includes two kinds: Model-based methods such as the SARSA algorithm and Model-irrelevant methods, like the Q-Learning algorithm.

III. REINFORCEMENT LEARNING

An agent learns to make decisions by interacting with the environment via a machine learning technique called reinforcement learning. Here there are unlabeled data, so the agent learns through experience. To maximise a reward signal, the agent engages with the surroundings. The agent learns the concept of a reward-based mechanism using the hit-and-trial method. It is employed in a variety of applications, including gaming, robotics, and decision-making.

A. Components of Reinforcement Learning

Agent- entity that interacts with the environment and makes decisions based on the feedback it receives. **Environment**- a world from which the agent operates and receives feedback. **State**- represents the environment at a given time, which the agent uses to make decisions. **State Space**- set of all possible states the agent could take to achieve the required goal. **Action**- decision that the agent makes based on the current state of the environment. **Action Space**- a Finite set of possible actions that the agent can take. **Reward**- Reward is a measurement metric that the environment sends back to the agent. **Policy**- strategy applied by an agent for the next action based on the current state. **Value**-Value is the anticipated long-term return at a discount in comparison to the short-term benefit. **Q-Value**-Similar to the value, Q-Value also requires a parameter called current action. It is the anticipated overall benefit of acting in a certain state and abiding by a certain set of rules.

B. Reinforcement Learning Algorithms

1) *Q-Learning*: A model-free reinforcement learning algorithm is called Q-Learning. Additionally, it has a reputation for being able to handle non-stationary situations where the

underlying distribution of the data may shift over time and capture complicated relationships among input features. In this research, temporal difference learning has been used. It is an algorithm that chooses the best action that should be taken at any possible state by estimating the value of each action known as a value-based algorithm. In Q-Learning, the agent maintains a Q-table which contains the estimated Q-values for each state-action pair. On each step, the agent selects an action to be performed based on its current state and updates the Q-table based on the observed reward and the estimated Q-values of the next state. Over time, as the agent experiences more states and updates its Q-table, the estimated Q-values converge to the true optimal Q-values, permitting the agent to learn the optimal policy. Thus, the algorithm guarantees to obtain optimal solutions.

Q-Learning guarantees global convergence to an optimal policy under certain conditions, which can provide more confidence in the learned policy than SARSA or DQL. This makes Q-Learning a more reliable option for SQL injection detection, where the consequences of false positives and false negatives can be severe.

2) *State Action Reward State action (SARSA)*: The main distinction between SARSA algorithms and Q-Learning is that SARSA algorithms do not require the greatest reward for the upcoming state in order to update the Q-value in the table. The action that the agent takes would be to either label the query as legitimate or malicious. The reward signal would be based on the accuracy of the agent's labelling. The agent employs SARSA during training to develop a policy that maximises the projected future reward. The agent starts in an initial state and selects an action based on its policy. It then observes the reward and the new state resulting from the action and updates its Q-value based on the SARSA update rule. Once the agent has learned a policy through training, it can be used to detect SQL injection attacks in real-time. The agent would observe a new SQL query and use its policy to label the query as legitimate or malicious.

The SARSA algorithm is an on-policy alternative to Q-learning, which means it learns a policy while taking into account the current policy being used. This can be useful in situations where the agent needs to balance exploration and exploitation in a more controlled manner, and where the reward signal may be noisy or delayed. Unfortunately, SARSA is known to converge slower than the Q-learning algorithm which can be an issue when training the algorithm on large datasets of SQL queries. SARSA has several hyperparameters that need to be carefully tuned to achieve optimal performance. This can be a difficult and time-consuming task, especially when dealing with large datasets.

3) *Deep Q-Learning*: Deep neural networks and reinforcement learning are combined in this system, which allows for more complex representations of the state and action spaces. A neural network is used in the Q-learning variant known as "Deep Q-Learning" (DQL) to approximate the Q-function, rather than a lookup table. During training, the agent uses DQL to learn a policy that maximizes the expected future

reward. The agent observes a state i.e., an SQL query, and passes it through a neural network to estimate the Q-values for each action i.e., legitimate or malicious. The agent selects the action with the highest Q-value based on an epsilon-greedy policy. The agent then observes the reward and the new state resulting from the action and updates the neural network using the back-propagation algorithm. Once the agent has learned a policy through training, it can be used to detect SQL injection attacks in real-time. A new SQL query is seen by the agent, which then runs it through the neural network to determine the Q-values for each action and chooses the one with the greatest Q-value. The agent then labels the query as legitimate or malicious based on the selected action.

Compared to traditional Q-Learning, DQL can handle high-dimensional state spaces more effectively by approximating the Q-function using a neural network. However, DQL requires a large amount of computation, particularly when using large neural networks. This can be a significant disadvantage when training the algorithm on large datasets of SQL queries. DQL can also be susceptible to overfitting, particularly when the dataset is small or noisy. Overfitting can lead to poor generalization to new SQL queries, reducing the algorithm's ability to detect SQL injection attacks.

C. Implementation of Q-Learning

1) *Algorithm*: The Q-learning algorithm constructs a table of Q-values that indicate the anticipated rewards for performing a certain action in a specific condition. The expected reward that can be attained by doing the action in the state in question and then implementing the best course of action is represented by the Q-value for that state-action pair. The algorithm starts by initializing the Q-table with arbitrary values, and then iteratively updates the Q-values based on the rewards received during each trial. An action (a) is taken based on the current state and Q-table. The Q-learning algorithm keeps going through iterations until the Q-values reach their ideal values. By choosing the action for each state that has the highest Q-value, the best policy can be found.

2) *Integrating Apriori with Q-Learning*: In Reinforcement Learning, the Apriori algorithm can be used to identify patterns in the interactions between an agent and its environment. They help the agent optimize its decision-making and improve its overall performance. If the agent frequently takes a certain action when it is in a specific state, the Apriori algorithm identifies this pattern and suggests that the agent take this action in the future whenever it is in a similar state. This helps the agent avoid sub-optimal decisions and improve its overall reward.

In our model, we provided the CNN model as an apriori to the Q-Learning algorithm. The CNN model was chosen over any other model like Naive Bayes because of its better ability to recognize a pattern in the input data that indicates the presence of an SQL injection attack. The variation in SQL injections is subtle, such as changes in the order of parameters or the use of different operators. CNN models are robust to these variations, as they can learn to recognize patterns

regardless of their location in the input data. CNN models also have the ability to learn complex patterns that are composed of simpler patterns. And most importantly CNN models are computationally efficient and can be trained on large datasets. This is important for SQL injection detection, as the model needs to be trained on a large and diverse set of input data in order to generalize well to new and unseen attacks.

Apriori can be integrated into a Q-Learning model at various stages of the learning process depending on the requirements. Apriori can be integrated at the pre-processing stage when identifying frequent patterns or correlations in the data and transforming the data into a more suitable format for the Q-Learning algorithm is crucial. Apriori can be used in the exploration/exploitation trade-off stage to balance the need for exploration with the need for exploitation and ensure that the Q-Learning algorithm is able to learn effectively from the available data. Apriori can also be integrated into the Q-Value update stage where certain state-action pairs are prioritized over others to improve the accuracy and efficiency of the Q-learning algorithm.

Thus, we will be integrating the Apriori of the CNN model in the Q-Value update stage. The CNN model will provide robustness while providing apriori at the Q-value update stage, which will improve the accuracy. In a Q-Learning environment model, a step function is a function that defines the behaviour of an agent as it interacts with its environment. The step function is used to update the Q-Value estimates for each state-action pair based on the observed reward and the predicted reward for the next state. Integrating the Apriori algorithm into the step function of a Q-Learning model involves using the interpretations of the apriori (CNN) model to modify the Q-Value estimates and the algorithm of the CNN model added as apriori to the step function of a Q-Learning environment model remains similar to the standard algorithm of Q-Learning model. The only difference is in the definition of the step function.

To integrate apriori in the Q-Learning model,

- Define the Q-learning environment model and its functions of it like reset and step. Also, define the Q-learning agent model and its function like act.
- The step function will take the query as input and with the help of the apriori of the CNN model it will determine the reward.
- The act function will make the prediction based on the policy the Q-Learning model has created.

3) *Updating of Q-Table*: The Q-value for a state-action pair (s, a) is updated using the observed reward after taking action (a) in state s and the predicted future rewards. The update is done as follows:

$$Q(s, a) = Q(s, a) + \alpha * (r + \gamma * \max_{a'}(Q(s', a')) - Q(s, a))$$

where, $Q(s, a)$ is the current estimate of the Q-value for state-action pair (s, a).

α is the learning rate, which determines the extent to which new information overrides the old estimate of the Q-value.




Fig. 1. Effect of Learning Rate on Accuracy




Fig. 2. Effect of Discount Factor on Accuracy

r is the reward observed after taking action a in state s .

γ is the discount factor, which determines the importance of future rewards relative to the present rewards.

$Q(s', a')$ is the estimated Q-value for the next state s' and the best action a' in that state.

IV. EXPERIMENTAL OUTCOMES

The Learning Rate (L.R.) is a trade-off between the algorithm's stability and its ability to respond to changes in the environment. It is a scalar value that ranges between 0 and 1. If the learning rate is high, the agent quickly assimilates new information and responds rapidly to the changes in the environment. However, a high learning rate makes the algorithm more unstable, since the agent overreacts to noisy information as seen in Fig. 1. If the learning rate is low, the agent updates its estimates more slowly and is less sensitive to changes in the environment. However, a low learning rate can also make the algorithm slow to meet the optimal policy since it takes longer for the agent to learn from its experiences.

The Discount Factor (D.F.) influences the Q-Learning algorithm's balance between short-term and long-term rewards. It is a scalar value that ranges between 0 and 1. A high discount factor (close to 1) means that future rewards are considered to be very important, and the agent will prioritize them over immediate rewards as seen in Fig. 2. A low discount factor (close to 0) means that the agent only cares about immediate rewards and does not consider future rewards.




Fig. 3. Effect of Learning Rate on Testing and Validation Predictions




Fig. 4. Effect of Discount Factor on Testing and Validation Predictions

We have proved that a higher learning rate does not necessarily result in better accuracy for the testing and validation data. While the accuracy for testing and validation data was relatively consistent across the varying learning rates, the true positive and false positive rates varied significantly. For example, a learning rate of 0.1 had a lower true positive rate and a lower false positive rate than a learning rate of 0.9 for the testing data, while the opposite was true for the validation data. Exploring different combinations of learning rates and discount factors yields better performance.

Also, the proposed system resulted in significant accuracy for both testing and validation datasets. The highest accuracy achieved for testing data was 0.7714, while for validation data it was 0.7905, both achieved with a discount factor of 0.7.

V. CONCLUSION

Using a Q-Learning algorithm for detecting SQL injection attacks in Web Applications holds promise in achieving accurate and real-time detection while minimizing false positives and negatives. The algorithm can be trained on a dataset of SQL queries and their labels to learn a policy for detecting malicious queries. Further research in this area can lead to the development of more sophisticated and robust algorithms for securing web applications against SQL injection attacks.

To increase the accuracy of the Q-Learning model, it is suggested to use a learning rate "closer" to 0 and a discount factor "closer" to 1 as seen from Fig. 1 and Fig. 2 where the accuracy for validation data is greater than the accuracy for testing data in an ideal condition. This increases the accuracy from 66% to around 70% for learning from training data. As the model learns from the training data, further learning with testing and validation data can lead to an increased accuracy from 70% to 76%. Fig. 3 and Fig. 4 visualize the values of True Positives and False Positives after learning from testing and then validation data with the varying values of Learning Rate and Discount Factor respectively. With a learning rate of 0.1 and discount factor of 0.9, not only can better accuracy values be obtained, but the difference between the number of actually detected injections and false alarms can also be maximized.

REFERENCES

- [1] "OWASP Top Ten", <https://owasp.org/www-project-top-ten>
- [2] Halfond, W. G., Viegas, J., and Orso, A. A Classification of SQL-Injection Attacks and Countermeasures. In SSSE (2006).
- [3] Wei, K., Muthuprasanna, M., & Suraj Kothari. (2006). Preventing SQL injection attacks in stored procedures. Australian Software Engineering Conference (ASWEC'06).
- [4] Rai, A., Miraz, M. M. I., Das, D., Kaur, H., & Swati. (2021). SQL Injection: Classification and Prevention. 2021 2nd International Conference on Intelligent Engineering and Management (ICIEM).
- [5] Erdödi, L., Sommervoll, Å. Å., & Zennaro, F. M. (2021). Simulating SQL injection vulnerability exploitation using Q-learning reinforcement learning agents. Journal of Information Security and Applications, 61, 102903.
- [6] Verme, M. D., Sommervoll, Å. Å., Erdödi, L., Totaro, S., & Zennaro, F. M. (2021, November). SQL Injections and Reinforcement Learning: An Empirical Evaluation of the Role of Action Structure. In Nordic Conference on Secure IT Systems (pp. 95-113). Springer, Cham.
- [7] Tang, P., Qiu, W., Huang, Z., Lian, H., & Liu, G. (2020). Detection of SQL injection based on artificial neural network. Knowledge-Based Systems, 105528. doi:10.1016/j.knosys.2020.105528
- [8] Niculae, S., Dichiu, D., Yang, K., & Bäck, T. (2020). Automating penetration testing using reinforcement learning.
- [9] Hu, Z., Beuran, R., & Tan, Y. (2020). Automated Penetration Testing Using Deep Reinforcement Learning. 2020 IEEE European Symposium on Security and Privacy Workshops (EuroS&PW).
- [10] Zennaro, F. M., & Erdodi, L. (2020). Modeling penetration testing with reinforcement learning using capture-the-flag challenges: trade-offs between model-free learning and a priori knowledge. arXiv preprint arXiv:2005.12632.
- [11] Ghanem, M. C., & Chen, T. M. (2019). Reinforcement learning for efficient network penetration testing. Information, 11(1), 6.
- [12] John, A. (2015). SQL Injection Prevention by adaptive algorithm. IOSR journal of computer engineering, 17, 19-24.
- [13] Hanmanthu, B., Ram, B. R., & Niranjan, P. (2015). SQL Injection Attack prevention based on decision tree classification. 2015 IEEE 9th International Conference on Intelligent Systems and Control (ISCO).
- [14] Sadeghian, A., Zamani, M., & Abdullah, S. M. (2013). A Taxonomy of SQL Injection Attacks. 2013 International Conference on Informatics and Creative Multimedia.
- [15] Medhane, M. H. A. S. (2013). Efficient solution for SQL injection attack detection and prevention. International Journal of Soft Computing and Engineering (IJSCE), 3, 396-398.
- [16] Qiang, W., & Zhongli, Z. (2011). Reinforcement learning model, algorithms and its application. 2011 International Conference on Mechatronic Science, Electric Engineering and Computer (MEC).